

Name: Aishwarya Autkar

Roll No: COTA06

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

```
In [4]: data = pd.read_csv('Hr.csv')
```

```
In [5]: data.columns
```


```
Out[5]: Index([u'EmpNumber', u'Age', u'Gender', u'EducationBackground',
u'MaritalStatus', u'EmpDepartment', u'EmpJobRole',
u'BusinessTravelFrequency', u'DistanceFromHome', u'EmpEducationLevel',
u'EmpEnvironmentSatisfaction', u'EmpHourlyRate', u'EmpJobInvolvement',
u'EmpJobLevel', u'EmpJobSatisfaction', u'NumCompaniesWorked',
u'Overtime', u'EmpLastSalaryHikePercent',
u'EmpRelationshipSatisfaction', u'TotalWorkExperienceInYears',
u'TrainingTimesLastYear', u'EmpWorkLifeBalance',
u'ExperienceYearsAtThisCompany', u'ExperienceYearsInCurrentRole',
u'YearsSinceLastPromotion', u'YearsWithCurrManager', u'Attrition',
u'PerformanceRating'],
dtype='object')
```

```
In [6]: data.head()
```

```
Out[6]:
```

	EmpNumber	Age	Gender	EducationBackground	MaritalStatus	EmpDepartment	EmpJobRole
0	E1001000	32	Male	Marketing	Single	Sales	Sales Executive
1	E1001006	47	Male	Marketing	Single	Sales	Sales Executive
2	E1001007	40	Male	Life Sciences	Married	Sales	Sales Executive
3	E1001009	41	Male	Human Resources	Divorced	Human Resources	Manager
4	E1001010	60	Male	Marketing	Single	Sales	Sales Executive

5 rows × 28 columns



In [7]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1200 entries, 0 to 1199
Data columns (total 28 columns):
EmpNumber          1200 non-null object
Age                1200 non-null int64
Gender             1200 non-null object
EducationBackground 1200 non-null object
MaritalStatus      1200 non-null object
EmpDepartment      1200 non-null object
EmpJobRole          1200 non-null object
BusinessTravelFrequency 1200 non-null object
DistanceFromHome    1200 non-null int64
EmpEducationLevel   1200 non-null int64
EmpEnvironmentSatisfaction 1200 non-null int64
EmpHourlyRate       1200 non-null int64
EmpJobInvolvement   1200 non-null int64
EmpJobLevel         1200 non-null int64
EmpJobSatisfaction  1200 non-null int64
NumCompaniesWorked  1200 non-null int64
OverTime            1200 non-null object
EmpLastSalaryHikePercent 1200 non-null int64
EmpRelationshipSatisfaction 1200 non-null int64
TotalWorkExperienceInYears 1200 non-null int64
TrainingTimesLastYear 1200 non-null int64
EmpWorkLifeBalance  1200 non-null int64
ExperienceYearsAtThisCompany 1200 non-null int64
ExperienceYearsInCurrentRole 1200 non-null int64
YearsSinceLastPromotion 1200 non-null int64
YearsWithCurrManager 1200 non-null int64
Attrition           1200 non-null object
PerformanceRating    1200 non-null int64
dtypes: int64(19), object(9)
memory usage: 262.6+ KB
```


In [8]: dept = data.iloc[:,[5,27]].copy()
dept_per = dept.copy()

In [9]: dept_per.groupby(by='EmpDepartment')['PerformanceRating'].mean()

Out[9]: EmpDepartment
Data Science 3.050000
Development 3.085873
Finance 2.775510
Human Resources 2.925926
Research & Development 2.921283
Sales 2.860590
Name: PerformanceRating, dtype: float64

```
In [10]: plt.figure(figsize=(10,4.5))
sns.barplot(dept_per['EmpDepartment'],dept_per['PerformanceRating'])
```

Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x7f70544a2c90>



```
In [11]: dept_per.groupby(by='EmpDepartment')['PerformanceRating'].value_counts()
```

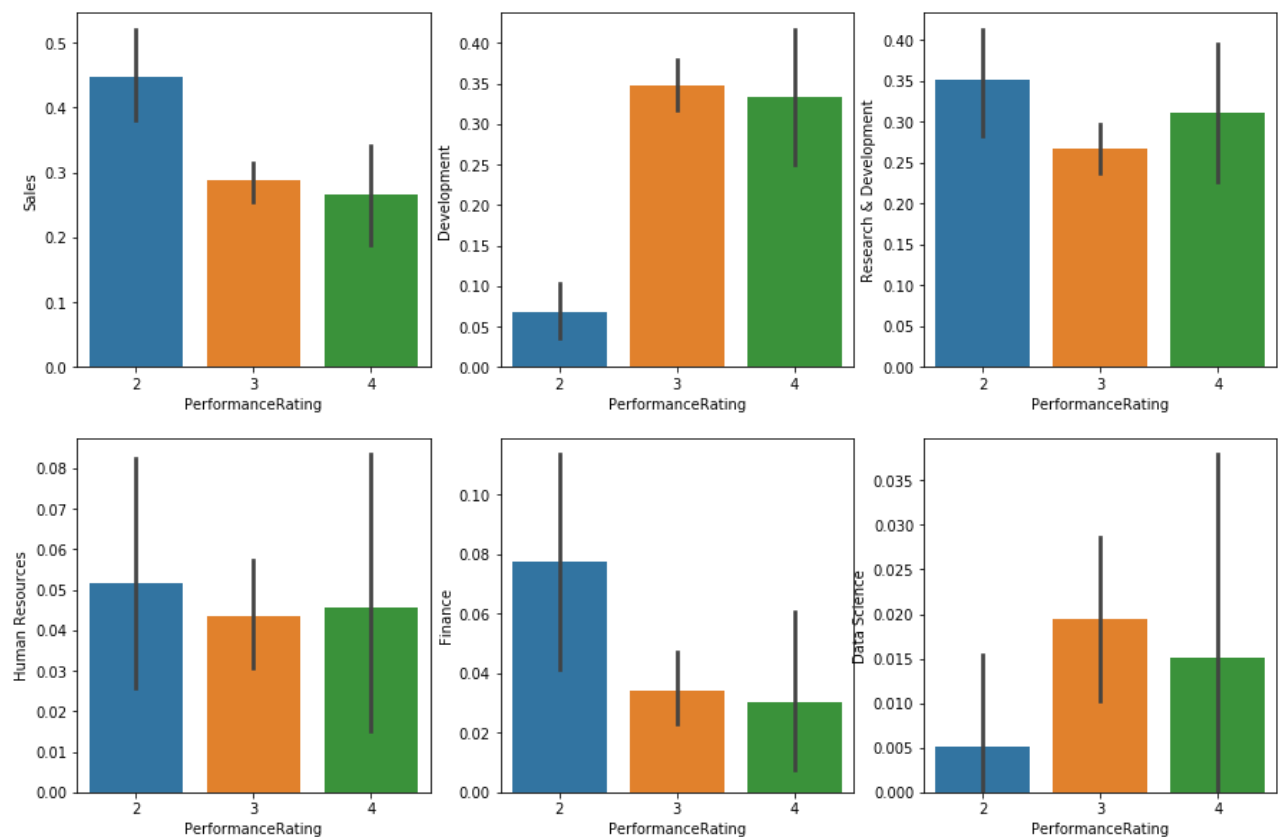
Out[11]:

EmpDepartment	PerformanceRating	
Data Science	3	17
	4	2
	2	1
Development	3	304
	4	44
	2	13
Finance	3	30
	2	15
	4	4
Human Resources	3	38
	2	10
	4	6
Research & Development	3	234
	2	68
	4	41
Sales	3	251
	2	87
	4	35

Name: PerformanceRating, dtype: int64

```
In [12]: department = pd.get_dummies(dept_per['EmpDepartment'])
performance = pd.DataFrame(dept_per['PerformanceRating'])
dept_rating = pd.concat([department,performance],axis=1)
```

```
In [13]: plt.figure(figsize=(15,10))
plt.subplot(2,3,1)
sns.barplot(dept_rating['PerformanceRating'],dept_rating['Sales'])
plt.subplot(2,3,2)
sns.barplot(dept_rating['PerformanceRating'],dept_rating['Development'])
plt.subplot(2,3,3)
sns.barplot(dept_rating['PerformanceRating'],dept_rating['Research & Development'])
plt.subplot(2,3,4)
sns.barplot(dept_rating['PerformanceRating'],dept_rating['Human Resources'])
plt.subplot(2,3,5)
sns.barplot(dept_rating['PerformanceRating'],dept_rating['Finance'])
plt.subplot(2,3,6)
sns.barplot(dept_rating['PerformanceRating'],dept_rating['Data Science'])
plt.show()
```



```
In [15]: enc = LabelEncoder()  
for i in (2,3,4,5,6,7,16,26):  
    data.iloc[:,i] = enc.fit_transform(data.iloc[:,i])  
data.head()
```

Out[15]:

	EmpNumber	Age	Gender	EducationBackground	MaritalStatus	EmpDepartment	EmpJobRc
0	E1001000	32	1	2	2	5	13
1	E1001006	47	1	2	2	5	13
2	E1001007	40	1	1	1	5	13
3	E1001009	41	1	0	0	3	8
4	E1001010	60	1	2	2	5	13


5 rows × 28 columns

In [16]: data.corr()

Out[16]:

	Age	Gender	EducationBackground	MaritalStatus	En
Age	1.000000	-0.040107	-0.055905	-0.098368	-0.
Gender	-0.040107	1.000000	0.009922	-0.042169	-0.
EducationBackground	-0.055905	0.009922	1.000000	-0.001097	-0.
MaritalStatus	-0.098368	-0.042169	-0.001097	1.000000	0.0
EmpDepartment	-0.000104	-0.010925	-0.026874	0.067272	1.0
EmpJobRole	-0.037665	0.011332	-0.012325	0.038023	0.5
BusinessTravelFrequency	0.040579	-0.043608	0.012382	0.028520	-0.
DistanceFromHome	0.020937	-0.001507	-0.013919	-0.019148	0.0
EmpEducationLevel	0.207313	-0.022960	-0.047978	0.026737	0.0
EmpEnvironmentSatisfaction	0.013814	0.000033	0.045028	-0.032467	-0.
EmpHourlyRate	0.062867	0.002218	-0.030234	-0.013540	0.0
EmpJobInvolvement	0.027216	0.010949	-0.025505	-0.043355	-0.
EmpJobLevel	0.509139	-0.050685	-0.056338	-0.087359	0.1
EmpJobSatisfaction	-0.002436	0.024680	-0.030977	0.044593	0.0
NumCompaniesWorked	0.284408	-0.036675	-0.032879	-0.030095	-0.
OverTime	0.051910	-0.038410	0.007046	-0.022833	-0.
EmpLastSalaryHikePercent	-0.006105	-0.005319	-0.009788	0.010128	-0.
EmpRelationshipSatisfaction	0.049749	0.030707	0.005652	0.026410	-0.
TotalWorkExperienceInYears	0.680886	-0.061055	-0.027929	-0.093537	0.0
TrainingTimesLastYear	-0.016053	-0.057654	0.051596	0.026045	0.0
EmpWorkLifeBalance	-0.019563	0.015793	0.022890	0.014154	0.0
ExperienceYearsAtThisCompany	0.318852	-0.030392	-0.009887	-0.075728	0.0
ExperienceYearsInCurrentRole	0.217163	-0.031823	-0.003215	-0.076663	0.0
YearsSinceLastPromotion	0.228199	-0.021575	0.014277	-0.052951	0.0
YearsWithCurrManager	0.205098	-0.036643	0.002767	-0.061908	0.0
Attrition	-0.189317	0.035758	0.027161	0.162969	0.0
PerformanceRating	-0.040164	-0.001780	0.005607	0.024172	-0.

27 rows × 27 columns




In [17]: data.drop(['EmpNumber'],inplace=True,axis=1)

In [18]: `data.head()`

Out[18]:

	Age	Gender	EducationBackground	MaritalStatus	EmpDepartment	EmpJobRole	Business1
0	32	1	2	2	5	13	2
1	47	1	2	2	5	13	2
2	40	1	1	1	5	13	1
3	41	1	0	0	3	8	2
4	60	1	2	2	5	13	2

5 rows × 27 columns




In [19]: `y = data.PerformanceRating`

In [20]: `X = data.iloc[:, [4,5,9,16,20,21,22,23,24]]`

In [21]: `X.head()`

Out[21]:

	EmpDepartment	EmpJobRole	EmpEnvironmentSatisfaction	EmpLastSalaryHikePercent	Em
0	5	13	4	12	2
1	5	13	4	12	3
2	5	13	4	21	3
3	3	8	2	15	2
4	5	13	1	14	3



In [22]: `X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3,random_state=10)`

In [23]: `sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)`

In [24]: `X_train.shape`

Out[24]: (840, 9)

In [25]: `X_test.shape`

Out[25]: (360, 9)

```
In [26]: # Training the model
from sklearn.ensemble import RandomForestClassifier
classifier_rfg=RandomForestClassifier(random_state=33,n_estimators=23)
parameters=[{'min_samples_split':[2,3,4,5],'criterion':['gini','entropy'],'min_samples_leaf':[1,2,3]}]
model_gridrf=GridSearchCV(estimator=classifier_rfg, param_grid=parameters, scoring='accuracy')
model_gridrf.fit(X_train,y_train)
```

```
Out[26]: GridSearchCV(cv=None, error_score='raise',
    estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
    max_depth=None, max_features='auto', max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=23, n_jobs=1,
    oob_score=False, random_state=33, verbose=0, warm_start=False),
    fit_params=None, iid=True, n_jobs=1,
    param_grid=[{'min_samples_split': [2, 3, 4, 5], 'criterion': ['gini', 'entropy'], 'min_samples_leaf': [1, 2, 3]}],
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring='accuracy', verbose=0)
```

```
In [27]: model_gridrf.best_params_
```

```
Out[27]: {'criterion': 'entropy', 'min_samples_leaf': 2, 'min_samples_split': 2}
```

```
In [28]: y_predict_rf = model_gridrf.predict(X_test)
```

```
In [29]: print(accuracy_score(y_test,y_predict_rf))
print(classification_report(y_test,y_predict_rf))
```

```
0.9305555555555556
```

	precision	recall	f1-score	support
2	0.92	0.89	0.90	63
3	0.94	0.97	0.96	264
4	0.83	0.73	0.77	33
avg / total	0.93	0.93	0.93	360

```
In [30]: confusion_matrix(y_test,y_predict_rf)
```

```
Out[30]: array([[ 56,   7,   0],
    [  4, 255,   5],
    [  1,   8,  24]])
```